


QWEN <-> V-JEPA semantic-axis alignment & visual groundedness across the cohort



The 3 group-QWEN PC axes are shared, interpretable semantic dimensions. ALIGNMENT (left, voxel-space): per axis, corr over reliable voxels of the QWEN- vs V-JEPA-band tuning (each = corr of a band's OUT-OF-FOLD per-video prediction with the PC-k scores). GROUNDEDNESS (right, brain-INDEPENDENT): CV ridge R^2 of predicting the axis from V-JEPA encoder features -- can a vision model recover a language-derived axis. Each subject individually (k-of-N).

WHAT WE EXPECT (and why): Concrete/seeable axes -> high alignment AND high groundedness in most subjects; a language-specific axis -> lower groundedness. Pre-registered order: dynamic/scene (PC2/PC1) >= social (PC0).

READING (candidate; the call is yours): RESULT: top-3 axis alignment & groundedness across 7 subjects.

