



QWEN <-> V-JEPA RSA across the cohort



QWEN<->V-JEPA representational alignment across subjects. Brain-independent FEATURE RSA (right) is the clean anchor -- do a language model and a vision model organize the videos alike, at all. Per-ROI model<->model similarity (left; from OUT-OF-FOLD predictions, ceiling-free) asks WHERE in cortex that shared semantic geometry is neural: it should CLIMB from early visual to category/STS cortex. Each subject individually (k-of-N).

WHAT WE EXPECT (and why): Shared semantic structure -> feature RSA > 0 in every subject and a per-ROI climb into category-selective cortex. Model<->brain magnitudes (not shown here) sit against an under-estimated 2-rep ceiling, so the model<->model curve is the robust cohort readout.

READING (candidate; the call is yours): RESULT: QWEN<->V-JEPA feature RSA + hierarchy across 7 subjects.